

11753 Computational Intelligence

Master in Intelligent Systems

Universitat de les Illes Balears

Handout #1: Feed-forward Neural Networks (FFNN)

T1. Follow the next indications to build an FFNN to solve the multi-class classification problem related to the *ic_lab1* dataset available at the course web page:

- (a) Load the *ic_lab1* dataset using the following source code:

```
1 import numpy as np
2 data = np.loadtxt('ic_lab1.txt')
3 X = data[:, :-1]
4 y = data[:, -1]
```

Listing 1: Loading of the *ic_lab1* dataset.

- (b) Determine the number of classes. For instance, you can use the following source code:

```
1 import numpy as np
2 M = len(np.unique(y))
```

Listing 2: Use of `numpy.unique` to find the number of classes.

- (c) Normalize the dataset and split it into the train and test sets using the function *train_test_split* of *scikit-learn* as follows:

```
1 from sklearn.model_selection import train_test_split
2 from sklearn.preprocessing import StandardScaler
3 scaler = StandardScaler()
4 X_ = scaler.fit_transform(X)
5 X_train, X_test, y_train, y_test = train_test_split(X_, y, test_size=.3, random_state=100)
```

Listing 3: Train and test splitting function.

for a 70% of the dataset devoted to training and 30% for testing (`random_state = <int value>` ensures a reproducible splitting [useful for debugging purposes]).

- (d) Make use of the indications in slides #4-6 of the lecture notes on *Keras* (& *Tensorflow*) to **define a neural network with one single hidden layer** with *nh* neurons. Let us assume that you store your network into the *model* object.

Remember that for a multi-class problem you need **as many output neurons as classes** *M*. (You can check diverse combinations of activation functions, as it is done in the source file *keras_examples.py* available in the course web page.)

- (e) **Define a training strategy**, i.e. an optimizer (e.g. *RMSprop*), a loss function (e.g. *categorical cross entropy*) and a metric (e.g. *accuracy*), in accordance to the indications of the slides #7-9, and **compile your model**. (You can check other, more complex strategies suggested in the source file *keras_examples.py* available in the course web page.)
- (f) Using slide #10 as a guide, **fit your model** using the training set, for a sufficient number of epochs, batch size and validation set. Remember that, previous to fitting a multi-class model, you need to transform your ground truth using *one-hot encoding*:

```
1 from keras.utils import to_categorical
2 y_train_ = to_categorical(y_train)
```

Listing 4: One-hot encoding Keras function.

- (g) Run your model for *nh* = 3 and *nh* = 4 and compare the resulting performances by means of suitable metrics and plots: loss and metric values, classification map and contour map. To this end, you can use slides #11-12 and the following source code (also available in *nn_evaluation.py* at the course web page):

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 from sklearn.metrics import confusion_matrix, accuracy_score
4 from keras.utils import to_categorical
```

```

5
6 def to_label(M, y):
7
8     if M == 2:
9         # round predictions
10        y_ = [round(x[0]) for x in y]
11    else:
12        # take max output
13        y_ = np.argmax(y, axis=1)
14
15    return y_
16
17 def plot_class(c, X, y):
18
19     i = np.where(y == c)[0]
20     plt.scatter(X[i,0],X[i,1])
21
22 def show_class_map_(model, X, y):
23
24     if X.shape[1] != 2:
25         print('This dataset is not two-dimensional ...')
26         return
27
28     M = len(np.unique(y))
29
30     # plot the contour map
31     plt.figure()
32     for c in range(M):
33         plot_class(c, X, y)
34     plt.axis('equal')
35
36     # plot the decision boundaries
37     ax = plt.gca()
38     xlim = ax.get_xlim()
39     ylim = ax.get_ylim()
40
41     # create the grid to evaluate the model at discrete feature space points
42     xx = np.linspace(xlim[0], xlim[1], 100)
43     yy = np.linspace(ylim[0], ylim[1], 100)
44     YY, XX = np.meshgrid(yy, xx)
45     xy = np.vstack([XX.ravel(), YY.ravel()]).T
46     ZZ_ = model.predict(xy)
47     ZZ = to_label(M, ZZ_)
48
49     # plot the boundaries
50     for c in range(M):
51         ax.contour(XX, YY, ZZ_[:,c].reshape(XX.shape), levels=[0.5], alpha=0.5, linestyle=['--'])
52
53     plt.xlabel('x1')
54     plt.ylabel('x2')
55     plt.axis('equal')
56     plt.title('contour map')
57     plt.show()
58
59     # plot the classification map
60     plt.figure()
61     plt.imshow(ZZ.reshape(XX.shape).T, origin='lower', extent=(xlim[0], xlim[1], ylim[0],
62         ylim[1]), cmap='RdYlGn')
63     plt.colorbar()
64     for c in range(M):
65         plot_class(c, X, y)
66     plt.xlabel('x1')
67     plt.ylabel('x2')
68     plt.axis('equal')
69     plt.title('classification map')
70     plt.show()
71
72 def do_show(model, history, X_train, y_train, X_test, y_test):
73
74     loss = history.history['loss']
75     val_loss = history.history['val_loss']
76     accuracy = history.history['accuracy']

```

```

76 val_accuracy = history.history['val_accuracy']
77
78 plt.figure()
79 plt.plot(loss, label='train')
80 plt.plot(val_loss, label='validation')
81 plt.title('loss function')
82 plt.legend()
83 plt.show(block=False)
84
85 plt.figure()
86 plt.plot(accuracy, label='train')
87 plt.plot(val_accuracy, label='validation')
88 plt.title('accuracy')
89 plt.legend()
90 plt.show(block=False)
91
92 M = len(np.unique(y_train))
93
94 # evaluation: train set
95 y_pred_ = model.predict(X_train)
96 y_pred = to_label(M, y_pred_)
97 cm = confusion_matrix(y_train, y_pred)
98 print(cm)
99 print('accuracy = %f' % (accuracy_score(y_train, y_pred)))
100
101 y_train_ = to_categorical(y_train)
102 score = model.evaluate(X_train, y_train_)
103 print("Train loss:", score[0])
104 print("Train accuracy:", score[1])
105
106 show_class_map_(model, X_train, y_train)
107
108 # evaluation: test set
109 y_pred_ = model.predict(X_test)
110 y_pred = to_label(M, y_pred_)
111 cm = confusion_matrix(y_test, y_pred)
112 print(cm)
113 print('accuracy = %f' % (accuracy_score(y_test, y_pred)))
114
115 y_test_ = to_categorical(y_test)
116 score = model.evaluate(X_test, y_test_)
117 print("Test loss:", score[0])
118 print("Test accuracy:", score[1])

```

Listing 5: Python functions to examine the performance of a neural network model. (Also available in *nn_evaluation.py*)

Although you should try to understand every line of the previous source code, pay special attention to:

- function *to_label*, at lines 6-15, which transforms the network output (= one neuron per class) into a numerical class label;
- line 51, which, for plotting the boundary for class *c* (against the other classes), we have to choose the output *c* of the prediction of the network for every grid point of the map *ZZ*_, and we select a level of 0.5 for drawing the boundary (= *levels* parameter);
- lines 101 and 115, which perform the one-hot encoding of the data prior to evaluating the network through the method *evaluate* of Keras models; and
- lines 95 and 109, which transform the network output into a numerical class label after a network prediction.

- T2. (a) Load the *iris* dataset and implement a solution for the inherent 3-class, 4-feature classification problem using the following **four neural network configurations**:
- one single hidden layer with 4, 7 and 10 neurons
 - two hidden layers with, respectively, 4 and 3 neurons
- (b) Use the following source code for evaluating the performance of each model using n-fold cross validation:

```

1 from sklearn.model_selection import KFold
2 def nfoldcv_keras(model, X, y):
3
4     M = len(np.unique(y))
5     n_folds = 5

```

```

6
7 results = np.zeros(n_folds)
8 kf = KFold(n_splits=n_folds, shuffle=True)
9 for k, (train, test) in enumerate(kf.split(X)):
10     # training
11     X_train, y_train = X[train], y[train]
12     y_train_ = to_categorical(y_train)
13     model.fit(X_train, y_train_, epochs=100, batch_size=10,
14               validation_split=0.2, verbose=0)
15     # testing
16     X_test, y_test = X[test], y[test]
17     y_pred_ = model.predict(X_test)
18     y_pred = to_label(M, y_pred_)
19     results[k] = accuracy_score(y_test, y_pred)
20
21 return results

```

Listing 6: n-fold cross validation of Keras models. (Also available in *nn_evaluation.py*)

You can next e.g. show a box-plot:

